

In [1]:

```
!pip install numpy  
!pip install pandas  
!pip install seaborn
```

Requirement already satisfied: numpy in c:\users\agrocel\anaconda3\lib\site-packages (1.24.3)  
Requirement already satisfied: pandas in c:\users\agrocel\anaconda3\lib\site-packages (2.0.3)  
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\agrocel\anaconda3\lib\site-packages (from pandas) (2.8.2)  
Requirement already satisfied: pytz>=2020.1 in c:\users\agrocel\anaconda3\lib\site-packages (from pandas) (2023.3.post1)  
Requirement already satisfied: tzdata>=2022.1 in c:\users\agrocel\anaconda3\lib\site-packages (from pandas) (2023.3)  
Requirement already satisfied: numpy>=1.21.0 in c:\users\agrocel\anaconda3\lib\site-packages (from pandas) (1.24.3)  
Requirement already satisfied: six>=1.5 in c:\users\agrocel\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas) (1.16.0)  
Requirement already satisfied: seaborn in c:\users\agrocel\anaconda3\lib\site-packages (0.12.2)  
Requirement already satisfied: numpy!=1.24.0,>=1.17 in c:\users\agrocel\anaconda3\lib\site-packages (from seaborn) (1.24.3)  
Requirement already satisfied: pandas>=0.25 in c:\users\agrocel\anaconda3\lib\site-packages (from seaborn) (2.0.3)  
Requirement already satisfied: matplotlib!=3.6.1,>=3.1 in c:\users\agrocel\anaconda3\lib\site-packages (from seaborn) (3.7.2)  
Requirement already satisfied: contourpy>=1.0.1 in c:\users\agrocel\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.0.5)  
Requirement already satisfied: cycler>=0.10 in c:\users\agrocel\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (0.11.0)  
Requirement already satisfied: fonttools>=4.22.0 in c:\users\agrocel\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (4.25.0)  
Requirement already satisfied: kiwisolver>=1.0.1 in c:\users\agrocel\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (1.4.4)  
Requirement already satisfied: packaging>=20.0 in c:\users\agrocel\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (23.1)  
Requirement already satisfied: pillow>=6.2.0 in c:\users\agrocel\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (9.4.0)  
Requirement already satisfied: pyparsing<3.1,>=2.3.1 in c:\users\agrocel\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (3.0.9)  
Requirement already satisfied: python-dateutil>=2.7 in c:\users\agrocel\anaconda3\lib\site-packages (from matplotlib!=3.6.1,>=3.1->seaborn) (2.8.2)  
Requirement already satisfied: pytz>=2020.1 in c:\users\agrocel\anaconda3\lib\site-packages (from pandas>=0.25->seaborn) (2023.3.post1)  
Requirement already satisfied: tzdata>=2022.1 in c:\users\agrocel\anaconda3\lib\site-packages (from pandas>=0.25->seaborn) (2023.3)  
Requirement already satisfied: six>=1.5 in c:\users\agrocel\anaconda3\lib\site-packages (from python-dateutil>=2.7->matplotlib!=3.6.1,>=3.1->seaborn) (1.16.0)

In [2]:

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

In [3]:

```
df = pd.read_excel(r'Agrocel_Jupyter.xlsx')
```

In [4]:

```
df
```

Out[4]:

	Date	Sulphur Rate
0	2021-04-01	23.31
1	2021-04-02	NaN
2	2021-04-03	22.96
3	2021-04-04	NaN
4	2021-04-05	NaN
...	...	...
960	2023-11-17	NaN
961	2023-11-18	NaN
962	2023-11-19	NaN
963	2023-11-20	NaN
964	2023-11-21	NaN

	Date	Sulphur Rate
0	2021-04-01	23.31
1	2021-04-02	NaN
2	2021-04-03	22.96
3	2021-04-04	NaN
4	2021-04-05	NaN
...	...	...
960	2023-11-17	NaN
961	2023-11-18	NaN
962	2023-11-19	NaN
963	2023-11-20	NaN
964	2023-11-21	NaN

965 rows × 2 columns

In [5]:

```
df.columns
```

Out[5]:

```
Index(['Date', 'Sulphur Rate'], dtype='object')
```

In [6]:

```
df = df.rename(columns={'Date': 'Date'})
```

In [7]:

```
from sklearn.impute import KNNImputer

# Initialize the KNN imputer
k = 5 # Number of neighbors to consider
imputer = KNNImputer(n_neighbors=k)

df_subset = df[['Sulphur Rate']]
df_imputed = imputer.fit_transform(df_subset)
df['Sulphur Rate'] = df_imputed
```

In [8]:

```
df['Date'] = pd.to_datetime(df['Date'])
```

In [9]:

```
df
```

Out[9]:

	Date	Sulphur Rate
0	2021-04-01	23.310000
1	2021-04-02	24.875247
2	2021-04-03	22.960000
3	2021-04-04	24.875247
4	2021-04-05	24.875247
...	...	...
960	2023-11-17	24.875247
961	2023-11-18	24.875247
962	2023-11-19	24.875247
963	2023-11-20	24.875247
964	2023-11-21	24.875247

	Date	Sulphur Rate
0	2021-04-01	23.310000
1	2021-04-02	24.875247
2	2021-04-03	22.960000
3	2021-04-04	24.875247
4	2021-04-05	24.875247
...	...	...
960	2023-11-17	24.875247
961	2023-11-18	24.875247
962	2023-11-19	24.875247
963	2023-11-20	24.875247
964	2023-11-21	24.875247

965 rows × 2 columns

In [10]:

```
print(df.dtypes)
```


Date	datetime64[ns]
Sulphur Rate	float64
dtype:	object

In [11]:

```
import matplotlib.pyplot as plt
import pandas as pd

# Convert 'Date' column to datetime
df['Date'] = pd.to_datetime(df['Date'])

# Plot the time series
plt.figure(figsize=(30, 20))
plt.plot(df['Date'], df['Sulphur Rate'], linestyle='-', color='b')
plt.title('Sulphur Rate Over Time', fontsize=22)
plt.xlabel('Date', fontsize=20)
plt.ylabel('Sulphur Rate', fontsize=20)
plt.grid(True)
plt.show()
```





In [12]:

```
from statsmodels.tsa.seasonal import seasonal_decompose  
  
result = seasonal_decompose(df['Sulphur Rate'],  
                           model ='multiplicative', period = 365)
```

In [13]:

```
result.plot()
```

Out[13]:



## ARIMA Model Preprocessing

In [14]:

```
#Train_test_split

spilt = (df.index < len(df)-30)
df_train = df[spilt].copy()
df_test = df[~spilt].copy()
```

In [15]:

```
#Stationarity checking

from statsmodels.tsa.stattools import adfuller

# Select the column 'Sulphur Rate' from the DataFrame
sulphur_rate_data = df_train['Sulphur Rate']

# Perform the Augmented Dickey-Fuller test
adf_test = adfuller(sulphur_rate_data)
print(f'p-value: {adf_test[1]}')      #This gives us a output of p-value: 0.6980113868350736
                                         #This result shows a large p-value, which means the tes
```

p-value: 0.6980113868350736


In [16]:

```
#First Differencing
```

```
df_train_diff = df_train['Sulphur Rate'].diff().dropna()  
df_train_diff.plot()
```

Out[16]:

<Axes: >



In [17]:


```
adf_test = adfuller(df_train_diff)  
print(f'p-value: {adf_test[1]}') # This gives us a p-value: 2.1229061985843574e-16 which is
```

p-value: 2.1229061985843574e-16

In [18]:


```
#Finding the parameters p and q
```

```
from matplotlib import pyplot
from statsmodels.graphics.tsaplots import plot_acf
series = df_train_diff #insert data here
plot_acf(series) #ACF plot function
pyplot.show() #Show graph
```



In [19]:

```
from matplotlib import pyplot
from statsmodels.graphics.tsaplots import plot_pacf
series = df_train_diff #insert data here
plot_pacf(series) #ACF plot function
pyplot.show() #Show graph
```



## Auto ARIMA

In [20]:

```
from pmdarima import auto_arima
```

```
model = auto_arima(df_train['Sulphur Rate'], seasonal=True, m=7, max_p=5, max_d=2, max_q=5,
print(model.summary())
```

### SARIMAX Results

```
=====
Dep. Variable:                      y      No. Observations:                  9
35
Model:                 SARIMAX(4, 1, 3)   Log Likelihood:           -2715.8
12
Date:                Tue, 19 Mar 2024   AIC:                         5447.6
24
Time:                    14:46:22     BIC:                         5486.3
40
Sample:                   0      HQIC:                         5462.3
87
                           - 935
Covariance Type:            opg
=====
5]
-----
-- ar.L1      -1.3429    0.028    -48.262    0.000    -1.397    -1.2
88
ar.L2      -1.6313    0.040    -40.603    0.000    -1.710    -1.5
53
ar.L3      -0.7510    0.039    -19.398    0.000    -0.827    -0.6
75
ar.L4      -0.3408    0.024    -14.244    0.000    -0.388    -0.2
94
ma.L1      0.0379    0.020     1.860    0.063    -0.002    0.0
78
ma.L2      0.1339    0.020     6.607    0.000     0.094    0.1
74
ma.L3      -0.8279    0.021    -39.526    0.000    -0.869    -0.7
87
sigma2     19.5688    0.708    27.626    0.000    18.180    20.9
57
=====
Ljung-Box (L1) (Q):                  0.15    Jarque-Bera (JB):
150.47
Prob(Q):                            0.70    Prob(JB):
0.00
Heteroskedasticity (H):              8.23    Skew:
-0.09
Prob(H) (two-sided):                0.00    Kurtosis:
4.96
=====
```

Warnings:

```
[1] Covariance matrix calculated using the outer product of gradients (complex-step).
```

In [32]:

```
from sklearn.metrics import mean_squared_error, r2_score

n_forecast = len(df_test) # Number of periods to forecast into the future
forecast = model.predict(n_periods=n_forecast)

# Step 9: Make Predictions
y_pred_test = forecast # No need to call .values

test_score = r2_score(df_test['Sulphur Rate'], y_pred_test)
print("Test Score (R-squared)(ARIMA Model):", test_score)

# Step 10: Evaluate the Model
mse = mean_squared_error(df_test['Sulphur Rate'], y_pred_test)
rmse = np.sqrt(mse)
print("Root Mean Squared Error(ARIMA Model):", rmse)

# Print the forecasts
print("Forecasted values:")
print(y_pred_test)
```

Test Score (R-squared)(ARIMA Model): -0.04121051848095836

Root Mean Squared Error(ARIMA Model): 5.062778070696675

Forecasted values:


```
935    23.276484
936    22.255916
937    21.396848
938    20.914060
939    20.534919
940    22.824571
941    21.023602
942    20.156218
943    22.668735
944    21.281856
945    20.310689
946    22.286083
947    21.402895
948    20.568364
949    21.977316
950    21.436691
951    20.791923
952    21.766029
953    21.435574
954    20.958698
955    21.626375
956    21.423894
957    21.077341
958    21.534145
959    21.410567
960    21.160578
961    21.472937
962    21.398416
963    21.218780
964    21.432201
dtype: float64
```

In [30]:

```
plt.figure(figsize=(10, 6))
plt.plot(df_test['Date'], df_test['Sulphur Rate'], label='Actual', color='blue')
plt.plot(df_test['Date'], y_pred_test, label='Predicted', color='red')
plt.title('Actual vs Predicted Sulphur Rate(Testing Data)')
plt.xlabel('Date')
plt.ylabel('Sulphur Rate')
plt.legend()
```

Out[30]:

&lt;matplotlib.legend.Legend at 0xcae80c15d0&gt;




In [29]:

```
forecast = model.predict(n_periods=len(df_train))
y_pred_train = forecast

plt.figure(figsize=(10, 6))
plt.plot(df_train['Date'], df_train['Sulphur Rate'], label='Actual', color='blue')
plt.plot(df_train['Date'], y_pred_train, label='Predicted', color='red')
plt.title('Actual vs Predicted Sulphur Rate(Training Data)')
plt.xlabel('Date')
plt.ylabel('Sulphur Rate')
plt.legend()
```

Out[29]:

&lt;matplotlib.legend.Legend at 0xcae8345c50&gt;



In [26]:

```
#Generate forecasts for the next 20 days
forecast_next_20_days = model.predict(n_periods=20)


# Print the forecasted values
print("Forecast for the next 20 days:")
print(forecast_next_20_days)
```

Forecast for the next 20 days:

```
935    23.276484
936    22.255916
937    21.396848
938    20.914060
939    20.534919
940    22.824571
941    21.023602
942    20.156218
943    22.668735
944    21.281856
945    20.310689
946    22.286083
947    21.402895
948    20.568364
949    21.977316
950    21.436691
951    20.791923
952    21.766029
953    21.435574
954    20.958698
dtype: float64
```

In [28]:

```
# Plot the forecasted values along with the existing data
plt.figure(figsize=(10, 6))
plt.plot(df['Date'], df['Sulphur Rate'], label='Historical Data', color='blue')
plt.plot(pd.date_range(df['Date'].iloc[-1], periods=20), forecast_next_20_days, label='Forecast')
plt.title('Sulphur Rate Forecast for the Next 20 Days')
plt.xlabel('Date')
plt.ylabel('Sulphur Rate')
plt.legend()
plt.grid(True)
plt.show()
```



## Manual ARIMA

In [34]:

```
from statsmodels.tsa.arima.model import ARIMA
model = ARIMA(df_train['Sulphur Rate'], order=(2,1,1))
model_fit = model.fit()
print(model_fit.summary())
```

SARIMAX Results

=====

==

Dep. Variable:	Sulphur Rate	No. Observations:	9
35			
Model:	ARIMA(2, 1, 1)	Log Likelihood	-2736.4
87			
Date:	Tue, 19 Mar 2024	AIC	5480.9
74			
Time:	16:46:08	BIC	5500.3
31			
Sample:	0	HQIC	5488.3
55			
Covariance Type:	opg		

=====

==

	coef	std err	z	P> z	[0.025	0.97
5]						
ar.L1	-0.4170	0.023	-17.760	0.000	-0.463	-0.3
71						
ar.L2	-0.2955	0.022	-13.615	0.000	-0.338	-0.2
53						
ma.L1	-0.8884	0.012	-76.439	0.000	-0.911	-0.8
66						
sigma2	20.4676	0.751	27.258	0.000	18.996	21.9
39						

=====

=====

Ljung-Box (L1) (Q):	0.75	Jarque-Bera (JB):
130.96		
Prob(Q):	0.39	Prob(JB):
0.00		
Heteroskedasticity (H):	9.17	Skew:
-0.04		
Prob(H) (two-sided):	0.00	Kurtosis:
4.83		

=====

=====

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

In [61]:

```
from sklearn.metrics import mean_squared_error, r2_score

# Step 8: Validate the Model
# Forecast
forecast = model_fit.forecast(steps=len(df_test))

# Step 9: Make Predictions
predictions = forecast.values

# Step 10: Evaluate the Model
mse = mean_squared_error(df_test['Sulphur Rate'], predictions)
rmse = np.sqrt(mse)
print("Root Mean Squared Error:", rmse)
```

Root Mean Squared Error: 4.929557193347942


In [62]:

In [63]:

```
plt.figure(figsize=(10, 6))
plt.plot(df_train['Date'], df_train['Sulphur Rate'], label='Actual', color='blue')
plt.plot(df_train['Date'], predictions, label='Predicted', color='red')
plt.title('Actual vs Predicted Sulphur Rate')
plt.xlabel('Date')
plt.ylabel('Sulphur Rate')
```

Out[63]:

Text(0, 0.5, 'Sulphur Rate')



In [64]:

```
from sklearn.metrics import r2_score

forecast = model_fit.forecast(steps=len(df_test))
predictions = forecast.values

# Assuming 'Sulphur Rate' is the target variable
y_true = df_test['Sulphur Rate'].values
y_pred = forecast.values

# Calculate R-squared score
test_score = r2_score(y_true, y_pred)
print("Test Score (R-squared):", test_score)

print("Predicted values:", y_pred)
print("True values:", y_true)
```


```
Test Score (R-squared): 0.012864920380618261
Predicted values: [23.41976753 22.69383408 20.18311815 21.4445569 21.660571
63 21.19770183
21.3268665 21.40980186 21.33704755 21.34287396 21.36194584 21.35227148
21.35066905 21.35419632 21.35319911 21.35257249 21.35312848 21.35308184
21.35293697 21.35301116 21.35302304 21.35299616 21.35300386 21.35300859
21.35300434 21.35300472 21.35300582 21.35300525 21.35300516 21.35300537]
True values: [24.87524725 24.87524725 24.87524725 16.64 24.87524725 1
3.9
24.87524725 24.87524725 13.9 24.87524725 13.9 24.87524725
24.87524725 13.9 24.87524725 13.67 24.87524725 13.67
24.87524725 24.87524725 24.87524725 24.87524725 24.87524725 13.67
13.67 24.87524725 24.87524725 24.87524725 24.87524725 24.87524725]
```

In [65]:

```
plt.figure(figsize=(10, 6))
plt.plot(df_test['Date'], df_test['Sulphur Rate'], label='Actual', color='blue')
plt.plot(df_test['Date'], predictions, label='Predicted', color='red')
plt.title('Actual vs Predicted Sulphur Rate')
plt.xlabel('Date')
plt.ylabel('Sulphur Rate')
```

Out[65]:

Text(0, 0.5, 'Sulphur Rate')



In [66]:

```
forecast_next_20_days = model_fit.forecast(steps=20)

# Print the forecasted values
print("Forecast for the next 20 days:")
print(forecast_next_20_days)
```


Forecast for the next 20 days:

935 23.419768  
936 22.693834  
937 20.183118  
938 21.444557  
939 21.660572  
940 21.197702  
941 21.326866  
942 21.409802  
943 21.337048  
944 21.342874  
945 21.361946  
946 21.352271  
947 21.350669  
948 21.354196  
949 21.353199  
950 21.352572  
951 21.353128  
952 21.353082  
953 21.352937  
954 21.353011

Name: predicted\_mean, dtype: float64

In [67]:

```
# Plot the forecasted values along with the existing data
plt.figure(figsize=(10, 6))
plt.plot(df['Date'], df['Sulphur Rate'], label='Historical Data', color='blue')
plt.plot(pd.date_range(df['Date'].iloc[-1], periods=20), forecast_next_20_days, label='Forecast')
plt.title('Sulphur Rate Forecast for the Next 5 Days')
plt.xlabel('Date')
plt.ylabel('Sulphur Rate')
plt.legend()
plt.grid(True)
plt.show()
```



In [ ]:

## VAR (Vector Autoregression Model)

In [38]:

```
df1 = pd.read_excel(r'sulphur and sulphuric acid daily data.xlsx')
```

In [39]:

df1

Out[39]:

	Date	Sulphur Rate	Sulphuric acid Rate
0	2021-04-01	23.31	NaN
1	2021-04-02	NaN	10.07
2	2021-04-03	22.96	10.07
3	2021-04-04	NaN	10.07
4	2021-04-05	NaN	10.07
...	...	...	...
960	2023-11-17	NaN	NaN
961	2023-11-18	NaN	NaN
962	2023-11-19	NaN	NaN
963	2023-11-20	NaN	NaN
964	2023-11-21	NaN	NaN

965 rows × 3 columns

In [40]:

df1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 965 entries, 0 to 964
Data columns (total 3 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Date             965 non-null    datetime64[ns]
 1   Sulphur Rate     364 non-null    float64 
 2   Sulphuric acid Rate 398 non-null  float64 
dtypes: datetime64[ns](1), float64(2)
memory usage: 22.7 KB
```

In [41]:

```
from sklearn.impute import KNNImputer

# Initialize the KNN imputer
k = 5 # Number of neighbors to consider
imputer = KNNImputer(n_neighbors=k)

df_subset = df1[['Sulphur Rate']]
df_imputed_sulphur = imputer.fit_transform(df_subset)
df1['Sulphur Rate'] = df_imputed_sulphur
```

In [42]:

df1

Out[42]:

	Date	Sulphur Rate	Sulphuric acid Rate
0	2021-04-01	23.310000	NaN
1	2021-04-02	24.875247	10.07
2	2021-04-03	22.960000	10.07
3	2021-04-04	24.875247	10.07
4	2021-04-05	24.875247	10.07
...	...	...	...
960	2023-11-17	24.875247	NaN
961	2023-11-18	24.875247	NaN
962	2023-11-19	24.875247	NaN
963	2023-11-20	24.875247	NaN
964	2023-11-21	24.875247	NaN

	Date	Sulphur Rate	Sulphuric acid Rate
0	2021-04-01	23.310000	NaN
1	2021-04-02	24.875247	10.07
2	2021-04-03	22.960000	10.07
3	2021-04-04	24.875247	10.07
4	2021-04-05	24.875247	10.07
...	...	...	...
960	2023-11-17	24.875247	NaN
961	2023-11-18	24.875247	NaN
962	2023-11-19	24.875247	NaN
963	2023-11-20	24.875247	NaN
964	2023-11-21	24.875247	NaN

965 rows × 3 columns

In [43]:

```
# Initialize the KNN imputer
k = 30 # Number of neighbors to consider
imputer = KNNImputer(n_neighbors=k)

df_subset = df1[['Sulphuric acid Rate']]
df_imputed_sulphuric = imputer.fit_transform(df_subset)
df1['Sulphuric acid Rate'] = df_imputed_sulphuric
```

In [44]:

df1.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 965 entries, 0 to 964
Data columns (total 3 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Date             965 non-null    datetime64[ns]
 1   Sulphur Rate     965 non-null    float64 
 2   Sulphuric acid Rate 965 non-null    float64 
dtypes: datetime64[ns](1), float64(2)
memory usage: 22.7 KB
```

In [45]:

```
#Stationarity checking

from statsmodels.tsa.stattools import adfuller

# Select the column 'Sulphur Rate' from the DataFrame
sulphur_rate_data = df1['Sulphur Rate']

# Perform the Augmented Dickey-Fuller test
adf_test = adfuller(sulphur_rate_data)
print(f'p-value: {adf_test[1]}')      #This gives us a output of p-value: 0.6980113868350736
                                         #This result shows a Large p-value, which means the tes
```

p-value: 0.682887920692923

In [46]:

```
#Stationarity checking

from statsmodels.tsa.stattools import adfuller

# Select the column 'Sulphur acid Rate' from the DataFrame
sulphur_rate_data = df1['Sulphuric acid Rate']

# Perform the Augmented Dickey-Fuller test
adf_test = adfuller(sulphur_rate_data)
print(f'p-value: {adf_test[1]}')      #This gives us a output of p-value: 0.6980113868350736
                                         #This result shows a Large p-value, which means the tes
```

p-value: 0.0033519626650673554


In [47]:

#First Differencing

```
df_diff = df1['Sulphur Rate'].diff().dropna()  
df_diff.plot()
```

Out[47]:

&lt;Axes: &gt;



In [51]:

```
#Stationarity checking

from statsmodels.tsa.stattools import adfuller

# Select the column 'Sulphur Rate' from the DataFrame
sulphur_rate_data = df_diff

# Perform the Augmented Dickey-Fuller test
adf_test = adfuller(sulphur_rate_data)
print(f'p-value: {adf_test[1]}')      #This gives us a output of p-value: 0.6980113868350736
                                         #This result shows a large p-value, which means the tes
```

p-value: 2.123006347762479e-17

In [54]:

```
from sklearn.model_selection import train_test_split
from statsmodels.tsa.api import VAR

# Step 1: Check Stationarity
acid_rate_stationary = adfuller(df1['Sulphuric acid Rate'])[1] < 0.05

# Step 2: Make "Sulphur Rate" Data Stationary
if not acid_rate_stationary:
    df['Sulphur Rate Diff'] = df1['Sulphur Rate'].diff().dropna()

# Step 3: Split the Data into Training and Testing Sets
train_df, test_df = train_test_split(df1, test_size=0.2, shuffle=False) # Adjust test_size

# Step 4: Fit VAR Model on the Training Set
if acid_rate_stationary:
    model = VAR(train_df[['Sulphur Rate', 'Sulphuric acid Rate']])
else:
    model = VAR(train_df[['Sulphur Rate Diff', 'Sulphuric acid Rate']])

results = model.fit()
```

In [55]:

```
results.summary()
```

Out[55]:

Summary of Regression Results

```
=====
Model:           VAR
Method:          OLS
Date:   Tue, 19, Mar, 2024
Time:    16:54:19
-----
No. of Equations: 2.00000  BIC:        4.86708
Nobs:            771.000  HQIC:       4.84483
Log likelihood: -4044.32  FPE:        125.326
AIC:             4.83092  Det(Omega_mle): 124.356
-----
Results for equation Sulphur Rate
```

```
=====
prob               coefficient      std. error      t-stat
-----
const              20.916100     1.385496     15.096
0.000
L1.Sulphur Rate   0.175403     0.035664     4.918
0.000
L1.Sulphuric acid Rate 0.030350     0.098566     0.308
0.758
-----
prob               coefficient      std. error      t-stat
-----
const              6.279709     0.478278     13.130
0.000
L1.Sulphur Rate   0.015249     0.012311     1.239
0.215
L1.Sulphuric acid Rate 0.333257     0.034025     9.794
0.000
-----
Correlation matrix of residuals
```

	Sulphur Rate	Sulphuric acid Rate
Sulphur Rate	1.000000	-0.046015
Sulphuric acid Rate	-0.046015	1.000000

In [ ]:

